

# Increasing Automation in the Backporting of Linux Drivers Using Coccinelle

Luis R. Rodriguez  
Rutgers University/SUSE Labs  
mcgrof@winlab.rutgers.edu

mcgrof@suse.com, mcgrof@do-not-panic.com

Julia Lawall  
Inria/LIP6/UPMC/Sorbonne University  
Julia.Lawall@lip6.fr

**Abstract**—Software is continually evolving, to fix bugs and add new features. Industry users, however, often value stability, and thus are not always able to update their code base to the latest versions. This raises the need to selectively backport new features to older software versions. Traditionally, backporting has been done by cluttering the backported code with preprocessor directives, to replace behaviors that are unsupported in an earlier version by appropriate workarounds. This approach however involves writing a lot of error-prone backporting code, and results in implementations that are hard to read and maintain. We consider this issue in the context of the Linux kernel, for which older versions are in wide use. We present a new backporting strategy that relies on the use of a compatibility library and on code that is automatically generated using the program transformation tool Coccinelle. This approach reduces the amount of code that must be manually written, and thus can help the Linux kernel backporting effort scale.

**Index Terms**—Linux, backports, program transformation

## I. INTRODUCTION

Linux is an open-source operating system kernel that has been under development since 1991. Its reliability, customizability, and low cost has made it a popular choice for an operating system kernel across the computing spectrum, from smartphones and tablets based on the Android distribution, to desktops based on popular distributions such as Debian and Ubuntu, to supercomputers based on Enterprise Linux releases. The Linux kernel is evolving rapidly, with a major release roughly every 2.5 months. Between the recent major releases v3.11 and v3.16 (September 2013 - August 2014), on average, 8,700 lines of code were added, 3,880 lines of code were removed, and 1,900 lines of code were modified *every day*.<sup>1</sup> This rapid rate of change combined with frequent releases allows the Linux kernel to keep up to date with bug fixes, new functionalities, and new services, such as new CPU architectures, device drivers and filesystems.

### A. The dilemma for silicon manufacturers

While the rapid release cycle of Linux has many benefits, it can be problematic for certain classes of users. Some system integrators may have invested heavily in testing a specific release and may wish to avoid regressions due to a kernel

upgrade. Upgrading a kernel may also require experience to understand what features to enable, disable, or tune to meet existing deployment criteria. In the worst case, some systems may rely on components that have not yet been merged into the mainline Linux kernel. Such a dependence may make it impossible to upgrade the kernel without cooperation from the component vendor or a slew of partners that need to collaborate on developing a new productized image for a system. As an example, development for 802.11n AR9003 chipset support on the upstream ath9k device driver started on March 20, 2010 with an early version of the silicon, at which point the most recent major release of the Linux kernel was v2.6.32. One of the first products to ship with this driver was the Google Chrome CR48, using ChromeOS, which started selling in retail in May 2011. The latest kernel release at this point was v2.6.38, but ChromeOS was still based on the v2.6.32 kernel, the release it was originally developed for.

The reluctance of users in many markets to keep up to date with the latest kernel poses a dilemma for silicon manufacturers, who need to make available device drivers so that their devices can be used on Linux systems. One approach is to develop drivers explicitly for the kernel releases that their clients are currently using. However, even clients who value stability may eventually need to modernize. Doing so then incurs the burden of a full rewrite or port of each driver to the newly adopted kernel release. While the device itself may be unchanged, the kernel evolves frequently with new internal APIs to improve the performance, robustness or flexibility of the code. Pervasive *collateral evolutions* are then needed to update the driver with respect to these changes [?]. And even once the driver is successfully ported to a more modern kernel, the problem is not really solved. The result will only be usable by those who are currently using the same kernel release.

An alternative to targeting a device driver to a specific kernel version is *upstream-first* development, in which code is initially developed only for the latest major kernel release, and then is submitted for inclusion *upstream*, i.e., into the Linux git repository maintained by Linus Torvalds, allowing the device driver to be included in the coming major release. While achieving inclusion upstream can be a challenge for silicon manufacturers, due to the strict coding guidelines of the Linux kernel, once it is achieved, the developers at the silicon company that upstreamed the device driver can then benefit

<sup>1</sup><https://github.com/gregkh/kernel-development/blob/1b17a1f21f2b0b871419f29947b2960cb36cb00b/kernel-development.pdf?raw=true>

from help from the Linux community in reviewing changes to the device driver as the Linux kernel evolves [?]. The silicon manufacturer needs to contribute only one complete version of the code; since it is upstream it will then be part of all future kernel releases and all future Linux distributions. As time goes by the silicon manufacturers may wish to turn their attention towards newer silicon. At this time, engineers can then orphan maintenance for older device drivers, allowing any community developer to take their maintenance on.

## B. Why and how Linux is backported

The upstream-first model makes device drivers available for users of future kernels, but leaves out those who must remain with older releases. If a device driver is only supported upstream, a Linux distribution or system integrator has no option but to backport that device driver down to the kernel release of interest. Backporting strategies have typically consisted of augmenting each affected file with `#ifdefs` that handle what is required for each kernel release on each target file. As each file is augmented individually, there is no code sharing, even within a single Linux distribution's codebase. Files also become harder to read, as the original code is interspersed with new code flows required to support each kernel release.

Since 2007, the Linux kernel backports project has promoted an alternative backport strategy, with the goal of maximizing code sharing and minimizing disruption to the individual driver source files, and with the goal of enabling upstream-first development of new drivers and features by making backports available to everyone, regardless of the Linux distribution used. The main innovation was to move the changes required to backport each driver out of the individual driver files and into a *compat* library, providing a set of helper functions. Indeed, typically, for a given class of devices, the drivers use a similar set of API functions and coding strategies, and these functions and coding strategies can all be backported in the same way. Rather than distributing the handling for each older release in every driver file, all of these variations are encapsulated into a *compat* library function. These functions can be declared as static inlines when performance is needed, or as external functions to limit the increase in code size. This approach was used in the ath9k support for ChromeOS noted above. The ath9k device driver was extended with AR9003 family chipset support upstream, and this support was incorporated as part of the v2.6.38 release at the time of the release of the Google Chrome CR48. Support for the AR9003 family of chipsets on ath9k on ChromeOS however was provided and backported onto the ChromeOS v2.6.32 based kernel using the Linux backports *compat* library.

The use of a *compat* library can dramatically reduce the amount of code changes required to support a class of drivers. Nevertheless, some changes per driver are still required, amounting to *glue code*, to invoke the *compat* library functions and to e.g., modify type definitions, which cannot be encapsulated into a function definition. These changes must initially

be made manually in each supported file to create *patches*,<sup>2</sup> which are made available to users, and these patches need to be maintained as the kernel evolves. Making these changes and maintaining the resulting patches is tedious and error prone, and limits the number of drivers that the backports project can support. A solution was thus needed to ease the process introducing this glue code.

## C. Our contribution

In this paper, we report on a new methodology for backports adopted by the Linux backports project that combines the use of a *compat* library with the use of Coccinelle [?]. Coccinelle is a program matching and transformation tool for C code that has been specifically designed to meet the needs and requirements of Linux developers. Matches and transformations are described in terms of an extension of the patch notation with generic features, resulting in *semantic patches*. Unlike patches, which are restricted to specific positions in specific files, a Coccinelle semantic patch expresses a change in a generic way, allowing it to be applied across an entire code base and to adapt to minor changes in this code base, as the code base evolves over time. In the context of backports, we automate the integration of the glue code into each supported driver using Coccinelle. Now the glue code needs to be specified only once, and then this specification can be applied automatically to all supported drivers over multiple releases, thus further reducing the amount of maintenance work to support a backport.

The rest of this paper is organized as follows. Section II provides the background for this work, including the relevant aspects of the Linux development models, the history of the backports project, and the use of Coccinelle. Section III then presents a tour of backporting strategies in more detail, based on a simple example. Next, Section IV expands this tour to illustrate a more complex case study, in which the changes required are determined by driver-specific information. Section V then highlights some correctness and performance issues. Finally, we briefly consider related work, conclusions, and directions for future work. This paper will put emphasis on how each new backporting strategy has helped the backports project to grow and scale, and ultimately how each of these strategies has helped to shift the objective of the project away from simply backporting the Linux kernel, towards trying to automate the backport process as much as possible.

## II. BACKGROUND

We now briefly review the starting points of our work: the Linux kernel development model, the Linux backports project and the Coccinelle program transformation tool.

### A. The Linux kernel development model

As illustrated in Figure 1, the development of a Linux kernel release is carried out in four phases: the merge window, the release candidate evaluation, the major release, and the

<sup>2</sup>A patch is a document indicating the lines of added and removed code, in a format generated by the Unix command `diff`. A patch can be automatically applied to a file using the Unix command `patch` [?].




Fig. 1: Four phases of Linux development

maintenance of a stable kernel. The merge window begins immediately after the previous major release, and lasts for 1-2 weeks. During this period, maintainers who have accumulated new features and device drivers for their subsystems request that Linus Torvalds collect and integrate their changes. This period culminates in the release of the first release candidate (rc1) making the complete set of merged changes available for general testing. The release of rc1 initiates the release candidate evaluation period. The purpose of this phase is to find and fix regressions for new code merged since the last major kernel release. There may be 5-9 release candidates, one per week. New drivers may be integrated during the release candidate evaluation period. Next comes the major release, making public the new version. At this point, the merge window for the next release begins. In parallel with the merge window, and beyond, possibly for several years, the current release is maintained as a stable version. Stable releases only contain critical bug fixes; they may never contain new features or drivers. Bug fixes are always sent upstream first, and must be merged into Linus Torvalds' tree before being cherry picked or ported to older maintained stable releases.

The Linux kernel development model introduces some delay between when new features are developed and when they are integrated into a release candidate or major release. To compensate for this delay, the `linux-next`<sup>3</sup> tree is used to mimic the merge window on daily basis, by pulling from each subsystem tree. Every day the tree is reset to the latest major kernel release and then each subsystem tree is pulled. The `linux-next` tree can thus be used to track the latest development efforts on all subsystems on a daily basis.

### B. A brief history of the Linux kernel backports project

The Linux kernel backports project<sup>4</sup> was started in 2007 by Luis R. Rodriguez while at Rutgers University, originally to help backport the 802.11 subsystem and a series of 802.11 device drivers to a series of older kernel releases.<sup>5</sup> The project was originally referred to as *compat-wireless*, reflecting the initial target. Over the years, the project grew to support more device drivers and subsystems. In April 2012, at the Linux Collaboration summit in San Francisco it was decided to rebrand the project *compat-drivers*<sup>6</sup> when the project was

folded under the Linux Foundation backports working group.<sup>7</sup> The project now spearheads the Linux kernel backports effort. The last rebranding of the project happened in April 2013, when it took on the name *backports*, to distinguish it from the Linux kernel compat layer, which addresses 64-bit and 32-bit compatibility. The backports project is now led by three core developers: Hauke Mehrtens, Johannes Berg, Luis Rodriguez; two co-maintainers: Hauke and Luis; and is developed and supported by a series of contributors. It backports the subsystems Ethernet, Wireless, Bluetooth, NFC, ieee802154, Media, and Regulator.

The current goal of the backports project is to backport a slew of device drivers from the latest major kernel releases down to a series of supported stable kernel releases, at a minimum those listed on the main kernel website, `kernel.org`. Currently, 18 releases are supported. The project's master development branch always tracks `linux-next`, allowing it to track all the development trees. This ensures that at the end of each merge window, the state of the backports will be very close to the state of the first release candidate. At this point, the backports project creates a further branch that tracks the progress of the new release over the release candidate evaluation period, to the major release, and on to its lifetime as a stable kernel. The backports project thus makes three kinds of backports available: those derived from `linux-next`, those derived from the most recent release candidate if any, and those derived from recent stable kernels. A user may prefer a backport from a stable kernel to one from `linux-next` or from a release candidate, if one is available for the desired driver.

As shown in Figure 2, as of September 2014, the backports project backports almost 800 drivers. These are kept up to date with `linux-next` and the recent stable kernels each day, and are guaranteed to at least compile correctly. Ensuring this each day typically requires 2-6 iterations of test, refinements, and compiles for all supported versions. For this development, the backports project uses a 32-core system with 236 GiB of RAM. Code generation and compile tests are all run in memory. As measured by GNU `time`, a full compilation test of a release across all 18 supported kernel revisions takes approximately 22 minutes of real time, 744 minutes of user mode time and 83 minutes of kernel time. When the backports project began in 2007, it provided backporting support for drivers down to v2.6.25, first released in 2008. In order to scale to a wider range of drivers, however, the project now only supports kernels down to at most Linux v3.0, first released in 2011. The original motivation behind the project was to encourage silicon manufacturers to work upstream on the Linux kernel while providing them a solution for backporting their drivers automatically down to older releases. The framework is designed *only* for Linux upstream drivers; the associated license enforces that proprietary drivers cannot and should not use this framework.

<sup>3</sup>[git://git.kernel.org/pub/scm/linux/kernel/git/next/linux-next.git](https://git.kernel.org/pub/scm/linux/kernel/git/next/linux-next.git)

<sup>4</sup><https://backports.wiki.kernel.org>

<sup>5</sup>[git://git.kernel.org/pub/scm/linux/kernel/git/mcgrof/compat-wireless-2.6-old.git](https://git.kernel.org/pub/scm/linux/kernel/git/mcgrof/compat-wireless-2.6-old.git)

<sup>6</sup>[git://git.kernel.org/pub/scm/linux/kernel/git/mcgrof/compat-drivers-old.git](https://git.kernel.org/pub/scm/linux/kernel/git/mcgrof/compat-drivers-old.git)

<sup>7</sup>[http://lists.linuxfoundation.org/pipermail/lf\\_driver\\_backport/2012-August/001075.html](http://lists.linuxfoundation.org/pipermail/lf_driver_backport/2012-August/001075.html)




Fig. 2: Number of backported drivers, for kernels released since 2009

### C. Coccinelle

Coccinelle is a program matching and transformation engine for C code [?], [?]. It is released as open source, and is available in a number of Linux distributions. Coccinelle provides a scripting language, SmPL (*semantic patch language*), that allows patterns to be expressed as fragments of C code, and transformations to be expressed by annotating lines of code with `-`, for removal of the matched code, or `+`, for addition of the corresponding code. As such, SmPL specifications resemble *patches* [?]. Nevertheless, they are more robust than ordinary patches in that they are insensitive to comments and whitespace, and take into account some aspects of the code semantics such as control flow and type information. Thus, we refer to SmPL specifications as *semantic patches*.

As a simple example of the use of Coccinelle, we consider the problem of replacing all calls to a function named `one` by calls to a function named `two`, and adding a `NULL` argument. The semantic patch that makes this change is shown in Figure 3. This semantic patch consists of a single rule that makes the complete transformation. The rule is in two parts: the declaration of *metavariables* that can match any term of a specified type between the initial pair of `@@` (lines 1-3), followed by a *transformation specification* (lines 4-5). In this case, the only metavariable is `arg`, which is declared to match any expression. The transformation specification then removes the call to `one` with its argument, while at the same time binding the metavariable `arg` to this argument expression. It then constructs a new function call to the function `two` with arguments the current binding of `arg` and `NULL`. This semantic patch can be applied to an entire code base. A more detailed presentation of Coccinelle is available in previous work [?], [?] and at the Coccinelle website.<sup>8</sup>

Coccinelle was originally motivated by a study of how the Linux kernel evolves [?]. This study identified the problem of *collateral evolutions*, in which the interface of a library changes, and all clients of the library must be updated accordingly. Coccinelle was designed to help Linux developers make collateral evolutions faster and more reliably. The problem

<sup>8</sup><http://coccinelle.lip6.fr>

Fig. 3: A simple semantic patch

of implementing backports using Coccinelle is related to the problem of collateral evolutions. While collateral evolutions were envisioned as being needed to modernize software, backports must address changes in library interfaces to transport modern code to older versions.

## III. BACKPORTING STRATEGIES

We now review three backporting strategies: i) the common strategy of providing version-specific implementations using `#ifdefs`, ii) the strategy originally taken by the Linux backports project of factorizing these changes into a compatibility library, and iii) our extension of the latter using Coccinelle.

### A. Backporting with `#ifdefs`

Backporting a device driver typically consists of modifying the source code with `#ifdefs` to handle the different requirements of different kernel releases. This entails adding blocks of code that provide alternate implementations for various functionalities, for different ranges of kernel versions, according to which evolutions have occurred and which collateral evolutions must be performed to accommodate them.

As a running example, we consider an evolution that was introduced by Linux kernel commit d314774cf2 and that was first merged upstream in Linux v2.6.29. This evolution moved a series of callback functions from the `net_device` data structure out into a new separate data structure of type `net_device_ops`. Backporting over this evolution, for Linux kernel versions before Linux v2.6.29, requires collateral evolutions that *undo* this change. Specifically, the definition of the new `net_device_ops` structure and the initialization of the link from the `net_device` structure to this new structure must be restricted, using an `#ifdef`, to the versions starting with Linux v2.6.29. Earlier versions must initialize the appropriate fields in the `net_device` structure itself, from among the callback functions that the modern driver puts in the new structure. Figure 4 shows a patch that backports this single collateral evolution on one device driver.

Lines 7 and 23 of Figure 4 introduce the `#ifdefs` that restrict some code of the modern driver to be used only in Linux versions v2.6.29 and later. Lines 25-31 introduce the code to be used for earlier versions, placing the relevant callback functions from the modern code (lines 8-13) into the fields of the `net` structure. In each case, the callback functions are used by the driver support library of the kernel, which is not backported, and which thus finds the desired functions in the expected place, with no further code changes.

In general, every driver that initializes a `net_device` structure would require all of these changes. Creating these

```

1 --- a/drivers/net/usb/usbnet.c
2 +++ b/drivers/net/usb/usbnet.c
3 @@ -1151,6 +1151,7 @@
4  }
5  EXPORT_SYMBOL_GPL(usbnet_disconnect);
6
7 +#if (LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,29))
8  static const struct net_device_ops usbnet_netdev_ops = {
9      .ndo_open = usbnet_open,
10     .ndo_stop = usbnet_stop,
11 @@ -1160,6 +1161,7 @@
12     .ndo_set_mac_address = eth_mac_addr,
13     .ndo_validate_addr = eth_validate_addr,
14 };
15 +#endif
16
17 /*-----*/
18
19 @@ -1229,7 +1231,15 @@
20     net->features |= NETIF_F_HIGHDMA;
21 #endif
22
23 +#if (LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,29))
24     net->netdev_ops = &usbnet_netdev_ops;
25 +#else
26 + net->change_mtu = usbnet_change_mtu;
27 + net->hard_start_xmit = usbnet_start_xmit;
28 + net->open = usbnet_open;
29 + net->stop = usbnet_stop;
30 + net->tx_timeout = usbnet_tx_timeout;
31 +#endif
32     net->watchdog_timeo = TX_TIMEOUT_JIFFIES;
33     net->ethtool_ops = &usbnet_ethtool_ops;

```

Fig. 4: Backporting the `net_device_ops` collateral evolution for the `usbnet` driver

patches, and maintaining them as other collateral evolutions are needed, is complex, tedious, and error prone.

### B. Backports via a compatibility library

Maintenance of patches is easy as long as the amount of changes being introduced is rather small. The `netdev_ops` collateral evolution, however, is an example of a collateral evolution that affects many network drivers, resulting in a large set of changes, that then have to be maintained in patch form. A better approach, proposed by the Linux backports project, consists of wrapping up the required changes into static inline or external helper functions and then using `#ifdefs` in these functions to adapt the code to each previous release.

This strategy is illustrated by the following code, which backports the `netdev_ops` collateral evolution for two device drivers. Now, the new `net_device_ops` structure used by the modern driver is maintained as is. Instead, we replace the direct initialization of the `netdev_ops` field by a call to a single static inline function defined by the backports compat library, amounting to glue code. Now, only this function needs multiple lines of `#ifdef` code, performing the direct assignment for versions starting with Linux v2.6.29, and copying the fields from the new structure into the main `net_device` structure for the older versions. Only one line of code is changed in each driver, in contrast to the 10 lines added to each driver by the previous approach.

```

1 --- a/drivers/net/usb/usbnet.c
2 +++ b/drivers/net/usb/usbnet.c
3 @@ -1446,7 +1446,7 @@ usbnet_probe (struct usb_interface *
4     udev

```

```

4     net->features |= NETIF_F_HIGHDMA;
5 #endif
6
7 - net->netdev_ops = &usbnet_netdev_ops;
8 + netdev_attach_ops(net, &usbnet_netdev_ops);
9     net->watchdog_timeo = TX_TIMEOUT_JIFFIES;
10    net->ethtool_ops = &usbnet_ethtool_ops;
11
12 --- a/drivers/net/wireless/ath/ath6kl/main.c
13 +++ b/drivers/net/wireless/ath/ath6kl/main.c
14 @@ -1289,7 +1289,7 @@ static const struct net_device_ops
15     ath6k
16
17 void init_netdev(struct net_device *dev)
18 {
19 - dev->netdev_ops = &ath6kl_netdev_ops;
20 + netdev_attach_ops(dev, &ath6kl_netdev_ops);
21     dev->destructor = free_netdev;
22     dev->watchdog_timeo = ATH6KL_TX_TIMEOUT;

```

Between 2007 and 2013 the backports project exclusively followed this strategy to help reduce the amount of maintenance on patches. The backport compat library now has a large set of helper functions that help to keep the number and size of the patches required for each backported driver to a minimum.

### C. The Coccinelle way to backport

The compat library approach reduces significantly the amount of code that must be modified in each driver. Still, backporting a new driver requires identifying the set of features that it uses, and comparing these features to those provided by the compat library to see where a collateral evolution to replace the existing code by glue code invoking the compat library is needed. Just as Coccinelle had been found to be useful in performing traditional (forward) collateral evolutions, we considered whether Coccinelle could also be useful for the kinds of collateral evolutions required in backports. For example, the `netdev_ops` collateral evolution could be expressed as a Coccinelle semantic patch as follows:

```

1 @@
2 struct net_device *dev;
3 struct net_device_ops ops;
4 @@
5 - dev->netdev_ops = &ops;
6 + netdev_attach_ops(dev, &ops);

```

This semantic patch could be used to backport the `netdev_ops` collateral evolution for all networking device drivers. It is indeed no longer even necessary for the developer to identify whether a new device driver to backport uses this features; Coccinelle both finds and updates the relevant code automatically. Note in particular that the semantic patch specifies the type of the expressions matching the `dev` and `ops` metavariables, to ensure that the transformation is performed only on structures of the appropriate type. Finally, this semantic patch amounts to only 6 lines of code to maintain, rather than 2 lines of code for each driver with the `#ifdef` approach.

## IV. CASE STUDY

To test the limits of what can be backported using Coccinelle, we chose the most complex collateral evolution supported by the backports project as a test case. Specifically, we decided to try to backport threaded IRQ support, introduced in the v2.6.31 kernel. This backport requires modifications to a driver-specific structure, as well as to multiple driver functions.

### A. Backporting threaded IRQ support the old way

We first explain how the backports project provided backport support for threaded IRQ support prior to the use of Coccinelle. The first step was to extend the compat library with support for threaded IRQs, as shown in Figure 5. This involved creating a new data type `compat_threaded_irq` (lines 1-11) to collect some extra information for each driver, and creating a set of helper functions to implement the threaded IRQ functionality (lines 13-69). The helper functions include `compat_request_threaded_irq` (lines 26-44), which initializes the fields of the `compat_threaded_irq` structure and then calls `request_irq`, and functions such as `compat_free_threaded_irq` (lines 46-51) and `compat_synchronize_threaded_irq` (lines 62-69) that call their unthreaded counterparts on information stored in the `compat_threaded_irq` structure. In all, the extension to the compat library amounts to 75 lines of code.<sup>9</sup>

Each driver to backport that relies on threaded IRQs then needs to be modified to make use of the new helper functions. Figure 6 shows the modifications for the b43 driver. Lines 1-13 update a header file to extend the driver’s private `b43_wldev` structure type with a field containing the compat structure when the kernel version is lower than the first one that supports threaded irq. Lines 14-52 replace each threaded IRQ operation with its compat version, again for kernels for which threaded irq. are not already supported.

The changes shown in Figure 6 amount to 20 lines of added code and only apply to a single driver. As of the linux-next of October 15, 2014, 169 files contain at least one call to `request_threaded_irq`, and of these 16 are in subsystems supported by the Linux backports project.<sup>10</sup> Backporting all of the 169 files that use threaded IRQs to Linux versions prior to v2.6.31 would require developing and maintaining over 3000 lines of patch code.

### B. Backporting threaded IRQ support with Coccinelle

Figure 7 shows a Coccinelle semantic patch that automates these changes. Most are relatively trivial: replace one call with another, with the new call using the backport data structure among its arguments. A typical example is illustrated in the first rule (lines 1-24), where a call to `request_threaded_irq` is replaced by a call to `compat_request_threaded_irq`. The new call takes the same arguments as the old one, with the addition of the first argument (line 17), which is the `compat_threaded_irq` structure.

A challenge in implementing this backport is where to store the `compat_threaded_irq` structure. As multiple instances of a device may be present in a running kernel, this structure cannot simply be a global variable of the device driver. In the case of the b43 driver, we placed this structure into the driver’s `b43_wldev` structure. To generalize this, we

```
1 #if LINUX_VERSION_CODE < KERNEL_VERSION(2,6,31)
2 struct compat_threaded_irq {
3     unsigned int irq;
4     irq_handler_t handler;
5     irq_handler_t thread_fn;
6     void *dev_id;
7     char wq_name[64];
8     struct workqueue_struct *wq;
9     struct work_struct work;
10 };
11 #endif
12
13 #if LINUX_VERSION_CODE < KERNEL_VERSION(2,6,31)
14 static inline
15 void compat_irq_work(struct work_struct *work)
16 {
17     ...
18 }
19
20 static inline
21 irqreturn_t compat_irq_dispatcher(int irq, void *dev_id)
22 {
23     ...
24 }
25
26 static inline
27 int compat_request_threaded_irq(
28     struct compat_threaded_irq *comp,
29     unsigned int irq,
30     irq_handler_t handler,
31     irq_handler_t thread_fn,
32     unsigned long flags,
33     const char *name,
34     void *dev_id)
35 {
36     comp->irq = irq;
37     comp->handler = handler;
38     comp->thread_fn = thread_fn;
39     comp->dev_id = dev_id;
40     INIT_WORK(&comp->work, compat_irq_work);
41     ...
42     return request_irq(irq, compat_irq_dispatcher, flags,
43         name, comp);
44 }
45
46 static inline
47 void compat_free_threaded_irq(
48     struct compat_threaded_irq *comp)
49 {
50     free_irq(comp->irq, comp);
51 }
52
53 static inline
54 void compat_destroy_threaded_irq(
55     struct compat_threaded_irq *comp)
56 {
57     if (comp->wq)
58         destroy_workqueue(comp->wq);
59     comp->wq = NULL;
60 }
61
62 static inline
63 void compat_synchronize_threaded_irq(
64     struct compat_threaded_irq *comp)
65 {
66     synchronize_irq(comp->irq);
67     cancel_work_sync(&comp->work);
68 }
69 #endif
```

Fig. 5: Extensions to the compat library to support `request_threaded_irq`

need to find a suitable location for this structure in each driver to which the semantic patch may be applied.

The need to support multiple instances of a data structure at run time is a common problem in device driver development,

<sup>9</sup>Computed using David Wheeler’s SLOCCount, <http://www.dwheeler.com/sloccount/>.

<sup>10</sup>1 in ethernet, 7 in wireless, 0 in bluetooth, 2 in nfc, 0 in ieee802145, 4 in media, 2 in regulator.



```

1 --- a/drivers/net/wireless/b43/b43.h
2 +++ b/drivers/net/wireless/b43/b43.h
3 @@ -805,6 +805,9 @@ enum {
4
5  /* Data structure for one wireless device (802.11 core) */
6  struct b43_wldev {
7  +#if LINUX_VERSION_CODE < KERNEL_VERSION(2,6,31)
8  + struct compat_threaded_irq irq_compat;
9  +#endif
10     struct b43_bus_dev *dev;
11     struct b43_wl *wl;
12     /* a completion event structure needed if this call is
13        asynchronous */
14 --- a/drivers/net/wireless/b43/main.c
15 +++ b/drivers/net/wireless/b43/main.c
16 @@ -4243,8 +4243,17 @@ redo:
17     if (b43_bus_host_is_sdio(dev->dev)) {
18         b43_sdio_free_irq(dev);
19     } else {
20 +#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,31)
21         synchronize_irq(dev->dev->irq);
22 +#else
23 + compat_synchronize_threaded_irq(&dev->irq_compat);
24 +#endif
25 +#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,31)
26         free_irq(dev->dev->irq, dev);
27 +#else
28 + compat_free_threaded_irq(&dev->irq_compat);
29 + compat_destroy_threaded_irq(&dev->irq_compat);
30 +#endif
31     }
32     mutex_lock(&wl->mutex);
33     dev = wl->current_dev;
34 @@ -4290,9 +4299,17 @@ static int b43_wireless_core_start(
35     goto out;
36 }
37 } else {
38 +#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,31)
39     err = request_threaded_irq(dev->dev->irq,
40         b43_interrupt_handler,
41         b43_interrupt_thread_handler,
42         IRQF_SHARED, KBUILD_MODNAME, dev);
43 +#else
44 + err = compat_request_threaded_irq(&dev->irq_compat,
45 +     dev->dev->irq,
46 +     b43_interrupt_handler,
47 +     b43_interrupt_thread_handler,
48 +     IRQF_SHARED, KBUILD_MODNAME, dev);
49 +#endif
50     if (err) {
51         b43err(dev->wl, "Cannot request IRQ-%d\n",
52             dev->dev->irq);

```

Fig. 6: Backporting the b43 driver

and the Linux kernel proposes a standard solution, the use of a *private* data structure. An instance of this private structure is created when a device is initialized, and then the driver infrastructure makes this structure available to each driver callback function, much like the implicit “this” argument found in object-oriented languages. Normally, each driver defines a specific private-structure type, containing the information that is specific to its operation. We exploit this private structure to store the `compat_threaded_irq` structure.

To use the private structure to store the `compat_threaded_irq` structure, we must address two issues. First, for each driver, we must find the type of the private structure and extend the corresponding type declaration with a field for the `compat_threaded_irq` structure. Second, we must find the name of the current instance of the private structure at each point where a reference to the `compat_threaded_irq` structure is needed for the backport process.

```

1 @ threaded_irq @
2 identifier ret;
3 expression irq, irq_handler, irq_thread_handler, flags,
4     name;
5 type T;
6 T *private;
7 @@
8
9 +#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,31)
10 ret = request_threaded_irq(irq,
11     irq_handler,
12     irq_thread_handler,
13     flags,
14     name,
15     private);
16 +#else
17 +ret = compat_request_threaded_irq(&private->irq_compat,
18 +     irq,
19 +     irq_handler,
20 +     irq_thread_handler,
21 +     flags,
22 +     name,
23 +     private);
24 +#endif
25
26 @ sync_irq depends on threaded_irq @
27 expression irq;
28 type threaded_irq.T;
29 T *threaded_irq.private;
30 @@
31
32 +#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,31)
33 synchronize_irq(irq);
34 +#else
35 +compat_synchronize_threaded_irq(&private->irq_compat);
36 +#endif
37
38 @ free_irq depends on threaded_irq @
39 expression irq, dev;
40 type threaded_irq.T;
41 T *threaded_irq.private;
42 @@
43
44 +#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,31)
45 free_irq(irq, dev);
46 +#else
47 +compat_free_threaded_irq(&private->irq_compat);
48 +compat_destroy_threaded_irq(&dev->irq_compat);
49 +#endif
50
51 @ modify_private_header depends on threaded_irq @
52 type threaded_irq.T;
53 @@
54
55 T {
56 +#if LINUX_VERSION_CODE < KERNEL_VERSION(2,6,31)
57 + struct compat_threaded_irq irq_compat;
58 +#endif
59 ...
60 };

```

Fig. 7: Backporting threaded IRQs with Coccinelle

To address the first issue, we exploit the fact that Coccinelle collects type information when analyzing the source code, and makes it possible to manipulate this type information via metavariables in a semantic patch. Fortunately, device drivers typically already pass their private structure as the last argument to `request_threaded_irq`, as the information contained in the private structure is typically also needed by the interrupt handler, which is installed by this function. By matching this reference to the private structure, Coccinelle makes it possible to obtain its type. Concretely, line 5 of Figure 7 declares a type metavariable `T`, which is

then used in describing the type of metavariable `private`. Matching `private` against the code in the last argument of `request_threaded_irq` has the side effect of storing the type of the matched code in `T`, where it can be used by subsequent rules. In the fourth rule (lines 51 to 60), `T`, referenced as `threaded_irq.T`, is used to match and extend the definition of the private structure, adding a new field `irq_compat` to the beginning of the private structure when the kernel version is less than v2.6.31.

To address the second issue, we exploit the fact that, within a given driver, the Linux developers typically give the private structure the same name, in every function in which it is used. Thus, we simply inherit the term matched by the metavariable `private` defined in the rule `threaded_irq`, and use that term in the added calls in the `synch_irq` and `free` rules. This solution is not safe, but it is pragmatic, in that it simplifies the transformation and exploits properties of the Linux coding style. Note that in the `free_irq` case (lines 38-49), we could also use the second argument to `free_irq`, which by definition of the IRQ API should point to the same structure as the last argument to `request_threaded_irq`.<sup>11</sup> This value is not immediately available in the `sync_irq` rule (lines 26-36), but we could extend the rule to match a neighboring expression of the right type, for a safer solution.

### C. Benefits of using Coccinelle for backporting

Reimplementing the threaded IRQ backport using Coccinelle revealed that the original manual backport was inconsistent. Specifically, in the manual backport, the `compat` structure, `compat_thread_irq`, was integrated into different structures in different drivers. Backporting is intrinsically risky, because the older code may not respect the invariants required by the backported code. Modifying the older code in a consistent way reduces the set of issues that can arise. Doing so also makes the results easier for developers to understand.

Reimplementing the threaded IRQ backport using Coccinelle also revealed that the existing threaded IRQ patch series also backported another collateral evolution, related to the management of IRQ flags. Isolating each collateral evolution in a separate patch benefits the backports project, by making it easier to understand how to backport other drivers, which may need only one of the changes. Using Coccinelle not only makes the need for this split apparent, it also makes it easier to manage the resulting set of changes. While two sets of changes are now needed, each is performed by a single semantic patch that can be applied to many files, rather than having to implement and record each of the changes individually.

Finally, for a few patch series, the amount of time to generate a backported release was actually reduced, as compared to the application of the manually created patches. Coccinelle has to parse the driver C code and search for relevant code, all of which are much more expensive than applying a patch, in which the file name, line number, and code

context are explicitly specified. Nevertheless, patch application is sequential, while Coccinelle can work on a collection of individual files in parallel. This difference is sufficient to make the use of Coccinelle faster in some cases.

## V. CORRECTNESS AND PERFORMANCE

We now consider some correctness and performance issues raised by the use of Coccinelle.

### A. Correctness

For the traditional uses of Coccinelle, for bug finding and collateral evolutions, a developer runs Coccinelle once on a code base, checks and possibly adjusts the results, and submits a patch upstream for review by kernel maintainers. The patch is integrated into the Linux kernel just like a manually generated patch, and Coccinelle is no longer involved.

The use case for backports is rather different. Here, the developer typically starts with a collection of patches, reflecting the result of backporting a number of drivers by hand. The developer then generalizes the existing patches into a semantic patch, which is then applied every day, as linux-next and the various stable kernels evolve. In this context, manually studying each result is not practical, and would likely not be reliable. Still, it is necessary to account for the possibility of errors, either in the semantic patch definition or in Coccinelle itself. Indeed, some improvements to Coccinelle were required to enable its use for backporting, such as improving the support for adding `#ifdefs` around complete function definitions.

To check the correctness of the result of Coccinelle, we exploit the existence of manually developed patches that perform the desired transformation. Specifically, we have developed a simple tool that applies the manually developed patches and the semantic patches to separate copies of the Linux kernel source code, and then checks that the results are equal. After applying the semantic patches to a possible target release and comparing the result with that of applying the manually written backport patches, we were surprised to find that the semantic patches produced more code changes. Indeed, the functionalities backported by the semantic patches occur in drivers that, while they are integrated on Linux backports, are not yet enabled for older kernels, as they require backporting of more collateral evolutions. A semantic patch transforms all code that exhibits the specified pattern, and thus generalizing a backport with Coccinelle can extend that backport to more drivers. Once all of the collateral evolutions required by a particular driver for a given kernel have been implemented, then that driver can be enabled for use on that kernel. When all required collateral evolutions to support a given kernel have been addressed with Coccinelle, then new drivers can be backported to that kernel fully automatically.

### B. Performance

Each day, the Linux kernel backports project generates and compile tests backport releases for all supported kernels, of which there are currently 18. Generation for all of them takes around 2 minutes on our 32-core machine, and compile

<sup>11</sup>Actually, unintentionally, in the second call that is added by this rule, this strategy is used.






Fig. 8: Generation time compared to number of patches, semantic patches, and supported kernel versions




Fig. 9: Number of backported drivers, for recent kernels

testing of the 18 resulting patched kernels takes 40 minutes of real time, comprising 1086 minutes of user time and 139 minutes of system time. Given the long compilation time, it is important to minimize the time for generating backports. Several solutions were explored to avoid the use of Coccinelle overwhelming the backporting time. Figures 8 and 9 relate the backport generation time to the number of supported kernels, patches, and resulting backported drivers for recent kernels.

As a first try, in v3.13, all of the semantic patches were concatenated and run in a single thread, much like the application of standard patches. In a uniprocessor environment, this avoids the cost of starting up Coccinelle for each semantic patch. This solution, however, does not exploit the parallelism inherent in applying a single specification to the hundreds of drivers supported by the backports project, and we see a high generation time in Figure 8 for v3.13. Coccinelle currently supports static parallelism via an external script that starts up  $n$  instances of Coccinelle and causes each of them to process  $1/n$  of the files in the code base. Performance can further be improved by instructing Coccinelle to use index information, precollected by a tool such as *glimpse*,<sup>12</sup> to identify the subset of files that may be relevant to a given semantic patch. With these two optimizations, semantic patch support can be less expensive than sequential application of the corresponding patches. For example, for v3.26.2, the number of supported drivers increases significantly (Figure 9), but the generation time actually slightly goes down. Currently, we are exploring the integration of parallelism into Coccinelle itself using the OCaml library *Parmap* [?], which provides dynamic scheduling. Preliminary tests suggest that this will improve CPU utilization and yield further performance improvements.

### C. Further improvements: a needle in a haystack

Coccinelle applies the rules of a semantic patch in order, from first to last, with any changes specified by a rule being

<sup>12</sup><http://webglimpse.net/>. As a side effect of this work, the first author convinced the developers of *glimpse* to release *glimpse* as open source.

performed as soon as the rule is matched. This can lead to extra matching and undesired modifications if rules that perform generic transformations are placed at the beginning of a semantic patch. As an example, consider the following rule that begins a semantic patch for backporting PCI drivers:

```
1 @ simple_dev_pm depends on module_pci @
2 identifier ops, pci_suspend, pci_resume;
3 declarer name SIMPLE_DEV_PM_OPS;
4 declarer name compat_pci_suspend;
5 declarer name compat_pci_resume;
6 @@
7 +compat_pci_suspend(pci_suspend);
8 +compat_pci_resume(pci_resume);
9 SIMPLE_DEV_PM_OPS(ops, pci_suspend, pci_resume);
```

This rule collects some information from each `SIMPLE_DEV_PM_OPS` structure, and introduces some PCI-specific calls from the `compat` library. The rule applies to every `SIMPLE_DEV_PM_OPS` declaration, and is not specific to PCI drivers in any way. Not only does it transform code that should not be transformed, it also can potentially have a significant performance impact. In the linux-next of October 15, 2014, there are 419 files that contain a `SIMPLE_DEV_PM_OPS` declaration, while only 61 of these are PCI drivers. All of these extra files will be parsed and (unnecessarily) transformed if the semantic patch is written in this way.

A solution is to add a rule that matches against some other more specific term that must be present if any transformation is to be performed. Other rules can then depend on the success of matching this rule. Essentially, this rule acts as a “needle in a haystack” to find the source files where the transformation is actually relevant. In this particular case, we observe that PCI drivers contain a `MODULE_DEVICE_TABLE` declaration with `pci` as the first argument. We thus add the following rule at the beginning of the semantic patch:

```
1 @ module_pci @
2 declarer name MODULE_DEVICE_TABLE;
3 identifier pci_ids;
4 @@
5 MODULE_DEVICE_TABLE(pci, pci_ids);
```

The `SIMPLE_DEV_PM_OPS` rule can now be declared to depend on the rule `module_pci`, ensuring that it is only applied to PCI drivers:

```
1 @ simple_dev_pm depends on module_pci @
2 identifier ops, pci_suspend, pci_resume;
3 declarer name SIMPLE_DEV_PM_OPS;
4 declarer name compat_pci_suspend;
5 declarer name compat_pci_resume;
6 @@
7 +compat_pci_suspend(pci_suspend);
8 +compat_pci_resume(pci_resume);
9 SIMPLE_DEV_PM_OPS(ops, pci_suspend, pci_resume);
```

The final rule to perform the backport uses metavariables that are defined by the previous one, and thus by that rule’s dependence will also only be applied to PCI drivers:

```
1 @@
2 identifier backport_driver;
3 expression pm_ops;
4 fresh identifier backports_pci_suspend = simple_dev_pm.
  pci_suspend ## "_compat";
5 fresh identifier backports_pci_resume = simple_dev_pm.
  pci_resume ## "_compat";
6 @@
7 struct pci_driver backport_driver = {
```

```

8 +#if (LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,29))
9     .driver.pm = pm_ops,
10 +#elif defined(CONFIG_PM_SLEEP)
11 + .suspend = backports_pci_suspend,
12 + .resume  = backports_pci_resume,
13 +#endif
14 };

```

## VI. RELATED WORK

The specific problem of backporting has not received much attention in the research community. Backporting is, however, related to more general issues of change management, as arise when merging trees in a source code management system and when integrating changes developed for one branch of a software product line into another branch.

Uquillas Gómez et al. [?] propose visualization tools to aid the developer in integrating a patch developed for one branch of a software project into another branch of the software project. They focus on individual changes and on helping the developer to identify semantic issues that may affect the correctness of the change in the new context. Other work on change impact analysis includes that of Gallagher and Lyle [?], who use program slicing [?] to collect information about the impact of a change, and the tool Chianti, which identifies change impact based on the results of test cases [?]. The work on change impact is complementary with ours. In our case, the correct backport is already identified, and we are concerned with expressing it in a concise and robust way. In the future, we could combine change impact analysis with our approach, to further check the correctness of the backported code.

Fiuczynski et al. faced the challenge of keeping an externally maintained patchset up to date with the evolutions in the Linux kernel [?]. They proposed a preliminary solution based on aspect-oriented programming [?] to re-express these patches in a more generic and robust way. To the best of our knowledge, this tool remained in a prototype stage.

## VII. CONCLUSIONS AND FUTURE WORK

The Linux backports project currently supports 5 semantic patches totaling 198 lines of code. These semantic patches affect 215 files in the linux-next of October 15, 2014. Recently, with v3.15, the Linux backports project has removed support for kernels older than v3.0 (see Figure 8). Six other semantic patches that had been developed for supporting older kernels were removed at this time. Furthermore, some glue code is still implemented by direct modification to the driver code. Using a semantic patch is appropriate when the changes are complex, are relevant to many drivers, and are susceptible to be affected by other evolutions in the code. Overall, the use of semantic patches has contributed to the robustness of the backport process. Indeed, another developer on the backports project recently stated “All the patches that broke often in the early days are now using coccinelle or are removed because they were only needed for the older kernel versions.”<sup>13</sup>

The work on backports raises a number of directions for future work. One direction would be to reduce the need for

glue code by integrating upstream static inline functions for accessing and updating key data structures. Patches to address this have been submitted and have now started to be accepted at least on the networking subsystem, maintained by David Miller, specifically to help reduce the amount of work to backport the ieee802154 subsystem.<sup>14</sup> this change is now merged as part of the v3.18-rc1 release.

Another direction would be to infer semantic patches. For many of our backports, we have a collection of manually written patches that make the same change. Backporting could be further streamlined by inferring semantic patches from these examples. Preliminary work has indeed been done on the automatic inference of change specifications [?], [?]. Alternatively, we observe that our additions of glue code amount to (the inverses of) collateral evolutions. If a library change is accompanied by a semantic patch, to ease updating the library’s clients, then it might be possible to systematically invert this semantic patch for subsequent backporting. Another direction would be to infer the glue code itself.

Finally, it is always a concern that a change in the code may break semantic invariants. We leave as future work to investigate whether change impact analysis, as described above, can be relevant here.

### A. License

This paper is licensed under the Creative Commons BY-SA 4.0.

<sup>13</sup>Hauke Mehrrens, private email of October 23, 2014.

<sup>14</sup><https://lkml.org/lkml/2014/4/17/663>